

Stay or Switch? Contributor-Aware Coding-Agent Reconfiguration After Repeated Non-Integration

Author Information Pending^{1*}

^{1*}Department Pending, Institution Pending, City Pending, Country Pending.

Corresponding author(s). E-mail(s): author@example.invalid;

Abstract

Context: Software repositories increasingly receive pull requests (PRs) produced with several coding-agent brands, yet cross-brand traces are often interpreted without separating tool change from contributor turnover. **Objective:** We study how repositories reconfigure coding-agent participation after an earlier PR outcome becomes known and after repeated non-integration. **Method:** Using a pinned AIDev revision containing 7,685,281 PRs, we construct 466,467 outcome-known, non-overlapping episodes from 31,323 repositories after observed multi-agent-brand onset. We classify episodes by agent-brand and contributor continuity, and construct exact three-predecessor decision chains. **Results:** An agent-brand change follows 31.2% of closed-unmerged episodes versus 14.2% of merged episodes; 83.3% of changes after closed-unmerged outcomes also change contributor identity. With contributor identity preserved, changing brand is associated with 72.7% 30-day integration versus 59.1% under persistence; the adjusted within-repository contrast is 7.88 percentage points (95% CI 2.08–13.68). Across exact non-integration streaks of one, two, and three-plus, observed switching declines from 26.9% to 19.0%, while the raw switch-minus-stay integration contrast rises from 2.04 to 17.98 points. **Conclusions:** Agent-brand change and contributor turnover are empirically distinct. The opposing history gradients reveal an adaptation gap but do not establish deliberate replacement, causal recovery, or agent-agent coordination.

Keywords: coding agents, pull requests, multi-agent ecosystems, software repositories, contributor turnover, mining software repositories

1 Introduction

Coding agents now participate directly in software repositories by opening pull requests, responding to review, and modifying evolving code bases. Architected systems such as ChatDev and MetaGPT organize several agents inside one designed workflow (Qian et al. 2023; Hong et al. 2023). Public repositories exhibit a different form of plurality: the same project may receive PRs labeled as Claude Code, Copilot, Cursor, Devin, Google Jules, or OpenAI Codex at different times. We call this *sequential multi-agent-brand participation*. It is an ecosystem-level observation and does not imply that agents communicate or coordinate.

Existing studies establish the scale of coding-agent activity (Li et al. 2026), compare agents and task types (Pinna et al. 2026), characterize non-integration (Ehsani et al. 2026; Peralta et al. 2026), and examine review and lifecycle roles (Nachuma and Zibran 2026; Chung and Hassan 2026). Concurrent cross-agent PRs and their merge conflicts have also been studied (Xu et al. 2026). Less is known about the sequential case in which an earlier PR is resolved before the next PR is opened. This ordering matters: it defines what outcome was observable at the next episode and avoids conflating sequential reconfiguration with concurrent work.

The central measurement problem is that changing an agent label is not necessarily a person changing tools. Contributor role and submitter attributes strongly shape agent use and PR integration (Cynthia et al. 2026; Yoshioka et al. 2026). If a different contributor opens the next PR, a change of agent brand can reflect team composition, task routing, or account-specific tool access. Treating every cross-brand episode as deliberate switching assigns intent that repository traces do not record.

We address this problem with a two-by-two transition taxonomy: same or different agent brand crossed with same or different contributor identity. Our unit is an outcome-known, non-overlapping pair of PRs in the same repository after the second observed agent brand appears. We ask three research questions:

RQ1 – Agent portfolios. How do repositories evolve from single-agent use to multi-agent-brand participation, and how often do observed brand transitions preserve contributor identity?

RQ2 – Outcome-conditioned reconfiguration. After an outcome-known merged versus closed-unmerged PR, how does the next episode redistribute across agent-brand continuity and contributor continuity?

RQ3 – Integration after non-integration. Following one or repeated closed-unmerged PRs, how are contributor-aware transition types associated with 30-day integration?

This paper makes three contributions. First, it defines an as-of-linked episode design in which every prior outcome is known before the next PR opens and prior episodes are not reused. Second, it separates observable agent-brand change from contributor turnover. Third, it documents an outcome-conditioned crossover and its limits: the association is not a general agent-performance ranking, is heterogeneous across repository strata, and does not license a causal recovery claim.

2 Background and Related Work	093
2.1 Coding-agent activity and integration	094
AIDev provides a multi-agent GitHub corpus and documents adoption and aggregate PR outcomes (Li et al. 2026). Studies using this ecosystem show why a global best-agent interpretation is unsafe. Task-stratified comparisons indicate that task mix can explain acceptance differences (Pinna et al. 2026). Analyses of merged and unmerged PRs identify change size, continuous integration, task, and review correlates (Ehsani et al. 2026; Nachuma and Zibran 2026). A qualitative analysis reports that only part of rejected agentic PRs can be clearly attributed to agent failure; workflow constraints and absent public rationale account for many cases (Peralta et al. 2026). We therefore use <i>closed-unmerged</i> as a neutral observed outcome.	095
2.2 Contributor and review context	096
Coding-agent PRs remain embedded in human governance. Initiation, iteration, approval, and merge roles differ across agent ecosystems (Chung and Hassan 2026). Core and peripheral developers delegate and review agent work differently (Cynthia et al. 2026); submitter attributes are also strong predictors of integration (Yoshioka et al. 2026). Reviewer exposure may change approval and scrutiny over time (Yu et al. 2026). These results motivate contributor continuity and calendar time as first-class parts of our design rather than afterthought controls.	097
2.3 Sequential versus concurrent multi-agent settings	098
Concurrent cross-agent PRs can edit overlapping files and exhibit elevated merge-conflict rates (Xu et al. 2026). Sequential benchmark work likewise shows that artifacts from one task can influence later tasks (Shastri et al. 2026). Our study differs from both: the prior PR is resolved before the current PR begins, and the estimand concerns observed repository reconfiguration rather than direct artifact interference. Work on algorithmic transference shows that people can generalize perceptions from one AI to another (Longoni et al. 2023), but our traces contain no measure of perception or trust. We use outcome carryover only as motivation, not as an observed psychological mechanism.	099
The targeted novelty update frozen on August 25, 2026 identified no direct study in the searched sources combining outcome-known ordering, non-overlap, agent-brand continuity, and contributor continuity. Because that update was a bounded scoping search rather than an exhaustive systematic review, we do not claim absolute priority.	100
3 Study Design	101
3.1 Dataset and analysis freeze	102
We use the Hugging Face dataset <i>AIDev-7.6M</i> (Li et al. 2026), pinned to revision 37bbe1533e26cc1e1374917dba1186d1c8a4dc81. The all-PR table contains 7,685,281 PRs	103

Table 1 Contributor-aware transition taxonomy.
Interpretations describe observable trace patterns, not verified intent.

Contributor	Agent brand	Operational label
Same	Same	Persistence
Same	Different	Brand change, same contributor
Different	Same	Contributor change, stable agent
Different	Different	Joint reconfiguration

from 961,168 repositories and six observed agent labels. Of these, 44,822 repositories contain at least two labels. The observation cutoff is March 31, 2026 at 23:59:59 UTC.

We parse creation, closure, and merge timestamps as UTC. A resolved prior PR must have a non-null closure at or after creation; if merged, its merge timestamp must also be at or after creation and strictly before the current PR opens. Ties in prior closure time within a repository are excluded because choosing one would be arbitrary. To ensure a complete 30-day outcome window, current PRs must open by March 1, 2026 at 23:59:59 UTC.

3.2 Outcome-known, non-overlapping episodes

For each repository, we find the first creation time for every observed agent brand. The second order statistic defines observed multi-agent-brand onset. The primary cohort includes only current PRs opened at or after this onset.

We link each current PR to the latest uniquely resolved PR whose outcome was known before current creation. This is an as-of join, not a row lag: a PR still open at current creation cannot be the predecessor. We require the prior merge outcome, when present, to be known before current creation. We then keep the first eligible current PR for each prior PR. This prevents a single prior outcome from being reused across several current observations.

The resulting ledger contains 466,467 episodes from 31,323 repositories. It contains 49,733 episodes after a prior closed-unmerged outcome and 416,734 after a prior merged outcome. Automated invariants verify unique current and prior episode identifiers, strict temporal ordering, post-onset membership, and a complete 30-day follow-up window.

3.3 Measures

The prior outcome is *merged* or *closed-unmerged*. The current integration outcome equals one when a valid merge timestamp falls from current creation through 30 days later. It measures repository integration, not code quality, trust, or defect freedom.

Agent-brand change equals one when current and prior agent labels differ. Contributor continuity equals one when their recorded `user_id` values match. There are no null or non-positive user identifiers in the primary ledger. We combine these measures into four mutually exclusive cells:

The same-contributor/different-brand cell is the strongest observable proxy for an individual changing tools. It remains imperfect because account sharing, automation, and decision intent are not observed.

3.4 Analysis

For RQ1, we center PR sequence position at each repository’s second observed agent-brand entry and report integration over ten PRs before and after onset. This event curve is a diagnostic: onset is endogenous, so pre-trends prohibit a causal event-study interpretation.

For RQ2, we estimate agent-brand switch rates and the four-cell transition distribution separately after merged and closed-unmerged prior outcomes. Confidence intervals use heteroskedasticity-robust standard errors clustered by repository.

For RQ3, we first report cell-specific 30-day integration rates with repository-clustered intervals. We then fit a linear probability model after closed-unmerged prior outcomes. Persistence is the reference cell. The model includes current-agent, prior-agent, calendar-month, and log-transformed gap controls. All variables, including the outcome and controls, are demeaned by repository, yielding a repository fixed-effects estimator. Standard errors are clustered by repository. Coefficients are associational because task choice, contributor experience, repository policy, and unobserved selection can affect both transition type and integration.

We then build a stricter three-predecessor decision chain. For each link, the predecessor’s disposition time is the later of closure and merge, and this time must strictly precede the successor’s creation. Exact streak one requires the immediately preceding PR to be closed-unmerged and its predecessor to be merged; exact streak two adds one earlier closed-unmerged PR followed by a merged PR; three-plus requires at least three consecutive closed-unmerged predecessors. Episodes without enough observed history to distinguish these categories are marked left-censored and excluded from the main gradient. This analysis contains 44,692 exact-history episodes and retains 5,519 censored episodes separately.

We additionally stratify the simpler switch-versus-stay contrast by repository stars. This is a transportability diagnostic, not evidence that stars themselves moderate a causal effect. Analyses are implemented in Python with Polars, pandas, statsmodels, Matplotlib, and pytest. The executable notebook and scripts record the pinned dataset revision and rebuild every reported table and figure.

3.5 Use of generative AI

OpenAI Codex was used during this working-draft stage to assist with code generation, search-log organization, exploratory analysis, and language drafting. The human authors are responsible for verifying the data construction, results, citations, and final text before submission. No generative model is listed as an author.

Merge rate around first multi-agent participation

Resolved PRs; event 0 is the first PR from a second distinct agent




Fig. 1 Integration around the second observed coding-agent-brand entry. Sequence position is defined within each repository. Error bars are descriptive binomial intervals. The pre-trend makes the curve unsuitable for a causal event-study interpretation.

4 Results

4.1 RQ1: onset is preceded by a pronounced integration trough

Figure 1 shows a gradual decline in integration before the second observed agent brand appears: from 85.2% ten PRs before onset to 75.8% one PR before onset. Integration is 78.3% at entry and then rises to 81.5% one PR later. The pre-onset trend is substantial, which is inconsistent with a clean exogenous intervention. We therefore interpret onset as occurring around a period of changing repository conditions or selected difficult work, not as a cause of the subsequent rebound.

Among closed-unmerged episodes followed by an agent-brand change, only 2,596 of 15,535 (16.7%) preserve contributor identity; 12,939 (83.3%) also change the recorded contributor. Thus, generic agent switching mainly describes joint contributor/agent reconfiguration rather than a within-person brand change.




Fig. 2 Contributor-aware transition composition and 30-day integration. Panel A partitions the next episode by prior outcome. Panel B reports repository-clustered point estimates and 95% confidence intervals after a closed-unmerged prior PR. These are observational associations.

4.2 RQ2: prior outcomes predict different reconfiguration patterns

An agent-brand change follows 31.2% of prior closed-unmerged episodes (repository-clustered 95% CI 29.0–33.5; $n = 49,733$), compared with 14.2% after prior merged episodes (95% CI 13.6–14.7; $n = 416,734$).

Panel A of Figure 2 decomposes this difference. After a closed-unmerged outcome, 63.1% of next episodes preserve both contributor and agent, 5.2% preserve the contributor but change the agent, 5.6% change the contributor but preserve the agent, and 26.0% change both. After a merged outcome, persistence rises to 83.6%, while joint reconfiguration falls to 10.7%. The behavioral contrast is therefore broader than agent choice: prior integration outcome is associated with both team and tool continuity.

4.3 RQ3: contributor-aware cells reveal opposing selection patterns

After a closed-unmerged prior PR, the current 30-day integration rate is 59.1% under persistence. It is 72.7% when the same contributor changes agent brand, 51.9% when the contributor changes but the agent remains stable, and 61.5% under joint reconfiguration (Figure 2, Panel B). The raw same-contributor brand-change contrast is therefore 13.63 percentage points.

In the repository fixed-effects model, same-contributor brand change remains associated with 7.88 percentage points higher integration relative to persistence (95% CI 2.08–13.68). Joint reconfiguration has a smaller adjusted association of 3.28 points (95% CI 1.64–4.91). Contributor change with a stable agent reverses from a negative




Fig. 3 Repository-demeaned switching contrasts from the aggregate robustness model. The model adjusts for current and prior agent, calendar month, and the inter-episode gap; intervals use repository-clustered standard errors.

raw contrast to a positive adjusted coefficient of 3.83 points (95% CI 1.20–6.46), indicating strong composition and timing differences. This reversal reinforces the need to avoid a mechanistic reading of unadjusted cells.

The earlier aggregate model provides a consistency check. Switching after a closed-unmerged outcome is associated with 6.48 points higher integration (95% CI 5.00–7.95), whereas switching after a merged outcome is associated with 4.29 points lower integration (95% CI –4.77 to –3.81) (Figure 3). This crossover is evidence against a universal benefit of switching.

Repository reputation provides a further boundary. The raw switch-minus-stay contrast after closed-unmerged outcomes is 3.55 points for zero-star repositories, 3.88 for 1–9 stars, 17.27 for 10–99 stars, and –1.51 for repositories with at least 100 stars; the last interval crosses zero (Figure 4). The full-corpus association therefore does not transport automatically to established projects.

4.4 RQ3 history: repeated non-integration exposes an adaptation gap

Figure 5 compares exact outcome-known histories. After a single closed-unmerged PR, 26.9% of next episodes change agent brand and the raw switch-minus-stay integration contrast is 2.04 points (repository-clustered 95% CI 0.44–3.64). After a streak of two, switching falls to 24.6% while the contrast rises to 6.43 points (95% CI 2.99–9.88). At three or more consecutive non-integrations, switching falls further to 19.0%, yet the contrast rises to 17.98 points (95% CI 6.97–28.99).

We call these opposing gradients an *adaptation gap*: observed agent-brand change becomes less common as repeated non-integration accumulates, although switch episodes are increasingly associated with integration relative to stay episodes. The pattern remains positive in both same- and changed-contributor strata, but these cells are compositionally imbalanced. The gradient is therefore a hypothesis-generating




Fig. 4 Raw switch-minus-stay 30-day integration contrasts after a closed-unmerged prior PR, stratified by repository stars. Intervals are descriptive normal-approximation intervals and do not adjust for selection.

boundary condition pending common-support weighting, not evidence that repeated non-integration causes switching to become effective.

5 Discussion

5.1 The empirical object is reconfiguration, not generic switching

The strongest descriptive finding is not that repositories simply try another agent after non-integration. Most observed brand changes also involve a new contributor. The two-by-two taxonomy changes the interpretation: same-contributor brand change is a small but analytically valuable cell; joint reconfiguration is far more common. This separation connects coding-agent research to established concerns about developer role and submitter attributes (Cynthia et al. 2026; Yoshioka et al. 2026).

The same-contributor association survives repository fixed effects and agent, month, and timing controls, making pure contributor replacement an insufficient explanation. It remains vulnerable to task routing and time-varying selection. A contributor may choose another agent precisely when the next task is easier, better specified, or better matched to that agent. Conversely, persistence may occur for unusually difficult continuations. The estimate should therefore motivate a task- and history-aware study, not a recommendation to switch tools.

The adaptation gap sharpens that agenda. Repeated non-integration may indicate harder tasks, stalled workflows, policy inertia, or repositories with fewer feasible alternatives. Any of these can produce decreasing switch propensity and increasing outcome

Agent switching across repeated non-integration

Exact outcome-known chains; repository-clustered 95% confidence intervals




Fig. 5 Observed switching and raw 30-day integration contrasts across exact, outcome-known non-integration streaks. Every predecessor disposition strictly precedes its successor’s creation; left-censored histories are excluded. Intervals use repository-clustered standard errors.

contrasts without a causal benefit. A confirmatory study should estimate overlap-weighted contrasts within streak and contributor strata, check directed agent-pair support, and audit whether consecutive PRs pursue the same underlying task.

5.2 Implications for research on agentic software engineering

First, evaluations should distinguish architected agent-agent collaboration from sequential multi-brand participation. The latter is observable at ecosystem scale but contains no direct coordination evidence. Second, provider comparisons should model directed substitutions and task mix instead of treating every brand change as homogeneous. Third, longitudinal analyses should enforce as-of knowledge. Ordering by row or creation time alone can attach an outcome that was not yet observable when the next PR began.

The null result in repositories with at least 100 stars is practically important. Mature projects may have stable automation policies, stronger review gates, or more deliberate task allocation. The apparent aggregate benefit is concentrated in less established repositories. Future work should use overlap weighting or matching to define a common-support population and should report transportability across repository strata.

5.3 Implications for repository maintainers and tool builders

Maintainers should not infer that a closed-unmerged PR proves agent failure; public rationale is frequently incomplete (Peralta et al. 2026). Nor should the results be used as a leaderboard. A safer operational implication is to retain structured provenance: agent brand, human initiator, task intent, review events, and reasons for closure. Such

metadata would allow teams to distinguish tool adaptation from contributor turnover and make later audits meaningful.

Tool builders can support this provenance by exposing stable machine-readable identifiers and recording handoffs. Repository interfaces could also surface whether an agent-authored PR continues an earlier attempt. These features would enable direct study of sequential recovery without inferring intent from labels.

6 Threats to Validity

Construct validity. Agent labels identify observed dataset categories, not necessarily the exact model, version, or interaction mode. Equal `user_id` is a proxy for contributor continuity and may fail under shared accounts or automation. Closed-unmerged is not verified failure, and merge within 30 days is not a quality measure. Multi-agent-brand participation does not imply agent-agent communication.

Internal validity. The study is observational. Task difficulty, contributor experience, repository policy, review availability, and unobserved workflow changes may determine both reconfiguration and merge. Repository fixed effects remove stable repository differences but not time-varying confounding. The strong RQ1 pre-trend and reversals between raw and adjusted contributor cells are direct warnings against causal interpretation.

Temporal validity. We use timestamps recorded by the dataset and exclude ambiguous ties and impossible orderings. Data errors, delayed closure, or later changes to GitHub state may remain. The 30-day cutoff excludes late merges by design. The corpus ends March 31, 2026, so later tools and policies are outside scope.

External validity. AIDev covers detectable public GitHub activity and does not represent private repositories or all uses of coding agents. The association differs sharply by repository stars and is null in the 100-plus-star stratum. Results should not be transported to mature projects, organizations, or other development platforms without replication.

Conclusion validity. Repository clustering addresses dependence within projects, while one-current-per-prior construction prevents repeated reuse of the same earlier outcome. However, multiple heterogeneity analyses are exploratory, task and review covariates are incomplete in the full corpus, and the current draft has not yet completed common-support weighting, directed agent-pair shrinkage, or a manual audit of sampled episode traces.

7 Conclusion

Sequential coding-agent traces show a clear outcome-conditioned reconfiguration pattern: closed-unmerged outcomes are followed by more brand change and more contributor turnover than merged outcomes. Yet the main measurement lesson is that these phenomena must be separated. Most cross-brand episodes also change contributor identity, while same-contributor brand change is less common and is associated with higher subsequent integration only in an observational, selection-sensitive design. The pattern is heterogeneous and absent in 100-plus-star repositories. The defensible

507 contribution is therefore a contributor-aware episode design and an empirical map of
508 reconfiguration, not a claim that switching causes recovery or that agents coordinate.

509

510 Acknowledgements

511

512 To be completed by the human authors before submission.

513

514 Declarations

515

516 **Funding.** Information pending author confirmation.

517

517 **Competing interests.** Information pending author confirmation.

518

518 **Ethics approval and consent to participate.** The study analyzes public repos-
519 itory metadata and does not recruit human participants. Institutional requirements
520 must be confirmed by the authors before submission.

521

521 **Consent for publication.** Not applicable.

522

522 **Data availability.** The source data are available from the [AIDev-7.6M dataset](#)
523 [page](#). This study pins revision 37bbe1533e26cc1e1374917dba1186d1c8a4dc81.
524 Redistribution and use remain subject to the source dataset and platform terms.

525

525 **Materials availability.** Search protocol, exact query log, and evidence map are
526 included in the companion artifact repository. The repository URL is pending archival
527 release.

528

528 **Code availability.** Analysis scripts, tests, an executed notebook, and derived
529 aggregate tables are prepared locally. A permanent public archive and identifier are
530 pending.

531

531 **Author contributions.** Information pending final authorship and CRediT role
532 confirmation.

533

534 References

535

536 Chung YJ, Hassan S (2026) Collaborator or assistant? how AI coding agents partition
537 work across pull request lifecycles. arXiv preprint arXiv:260508017 [https://doi.or](https://doi.org/10.48550/arXiv.2605.08017)
538 [g/10.48550/arXiv.2605.08017](https://doi.org/10.48550/arXiv.2605.08017)

539

540 Cynthia ST, Das JK, Roy B (2026) Are we all using agents the same way? an empirical
541 study of core and peripheral developers’ use of coding agents. In: Proceedings of
542 the 23rd International Conference on Mining Software Repositories, pp 520–532,
543 <https://doi.org/10.1145/3793302.3793377>

544

545 Ehsani R, Pathak S, Rawal S, et al (2026) Where do AI coding agents fail? an empirical
546 study of failed agentic pull requests in GitHub. arXiv preprint arXiv:260115195
547 <https://doi.org/10.48550/arXiv.2601.15195>

548

549 Hong S, Zhuge M, Chen J, et al (2023) MetaGPT: Meta programming for a multi-
550 agent collaborative framework. arXiv preprint arXiv:230800352 [https://doi.org/10](https://doi.org/10.48550/arXiv.2308.00352)
551 [.48550/arXiv.2308.00352](https://doi.org/10.48550/arXiv.2308.00352)

552

Li H, Zhang H, Hassan AE (2026) AIDev: Studying AI coding agents on GitHub. In: Proceedings of the 23rd International Conference on Mining Software Repositories, pp 1029–1033, https://doi.org/10.1145/3793302.3797249	553 554 555 556
Longoni C, Cian L, Kyung EJ (2023) Algorithmic transference: People overgeneralize failures of AI in the government. <i>Journal of Marketing Research</i> 60(1):170–188. https://doi.org/10.1177/00222437221110139	557 558 559 560
Nachuma C, Zibran M (2026) When AI teammates meet code review: Collaboration signals shaping the integration of agent-authored pull requests. arXiv preprint arXiv:260219441 https://doi.org/10.48550/arXiv.2602.19441	561 562 563 564
Peralta SRO, Hoshi F, Washizaki H, et al (2026) Why are agentic pull requests merged or rejected? an empirical study. arXiv preprint arXiv:260522534 https://doi.org/10.48550/arXiv.2605.22534	565 566 567 568
Pinna G, Gong J, Williams D, et al (2026) Comparing AI coding agents: A task-stratified analysis of pull request acceptance. arXiv preprint arXiv:260208915 https://doi.org/10.48550/arXiv.2602.08915	569 570 571 572
Qian C, Liu W, Liu H, et al (2023) ChatDev: Communicative agents for software development. arXiv preprint arXiv:230707924 https://doi.org/10.48550/arXiv.2307.07924	573 574 575 576
Shastri KNA, Senrayan G, Satapara S, et al (2026) Beyond isolated tasks: A framework for evaluating coding agents on sequential software evolution. arXiv preprint arXiv:260403035 https://doi.org/10.48550/arXiv.2604.03035	577 578 579 580
Xu G, Subramanian A, Karthik N (2026) AI agent pull requests on GitHub: Frequency, structure, and merge conflict rates. arXiv preprint arXiv:260704697 https://doi.org/10.48550/arXiv.2607.04697	581 582 583 584
Yoshioka H, Monno T, Tokumasu H, et al (2026) Let’s make every pull request meaningful: An empirical analysis of developer and agentic pull requests. In: Proceedings of the 23rd International Conference on Mining Software Repositories, pp 747–751, https://doi.org/10.1145/3793302.3793567	585 586 587 588
Yu H, Liu L, Jiang X, et al (2026) Habituation at the gate: Rising approval and declining scrutiny in human review of AI agent code. arXiv preprint arXiv:260622721 https://doi.org/10.48550/arXiv.2606.22721	589 590 591 592 593 594 595 596 597 598